

Lesson 3.3: Monitoring with a DevOps Mindset

DEVOPS CULTURE AND MINDSET

Monitoring with a DevOps Mindset



Courtney Kissler
Vice President
Digital Platform Engineering
Nike

UCDAVIS

Continuing and Professional Education



Learning Objectives

Explain why incident dashboards aren't always helpful

Discuss importance of transitioning to proactive monitoring

Identify business metrics for monitoring



Monitoring Has Evolved

Old way: Focus was primarily on system and infrastructure metrics

Related to **availability** and **performance**

“Is the server up?”

Performance measured inconsistently



Boy Who Cried Wolf DevOps

Reactive dashboards sounded the alarm

No proactive alerts before a problem

False alarms caused alerts to be ignored



Old Way of Monitoring was **NOT** a Recipe for Long-term Success


Server up but application not available

Sending alerts that required no action

Stable infrastructure but *unstable* apps

Slide 5: We Realized Our Monitoring Had to Evolve

We Realized Our Monitoring Had to Evolve





Business Metrics as a Way of Monitoring


Metric dip alerts team to investigate

Started tracking **M**ean **T**ime **T**o **D**etect

No more 30 minutes to calls to help desk!

Automated alert triggers system to solve in seconds

Slide 7: Accelerate





Observability is Becoming the Industry Norm

You have the instrumentation you need to understand what's happening within your software



Charity Majors on Benefits of Observability

Thought leader on current best practices

Approach to monitoring is lagging

Health of the **system** no longer matters

What matters is health of each event



Observability Instead of Monitoring

It's now about **unknown** unknowns

It's no longer about known unknowns

Gone are the days...


You Can't Monitor Everything

Engineers have vital **responsibilities**

Auto-remediate what you can

Fail gracefully where you can't

Shift ops load to providers as much as possible



Learn More from Charity Majors

You can't predict all **unknown** unknowns

Focus energy where you can **predict**:

- Instrumentation

- Resilience to failure

- Fast and safe deployment



Stop the Sustained Barrage

If you know how to fix it, **fix it**

Auto-remediate the problem

Disable paging alerts in off hours

Make **system resilient** enough to wait



Check Out Charity Majors' Company: Honeycomb

Explore this further

See amazing sampling documents

Check out more **best practices**



Summary

The **importance** of monitoring cannot be overstated

Transition to some form of proactive monitoring

Identify and **diagnose** unknown unknowns

Observe business metrics and sampling to provide early warnings